

AlgoLab Graph and BGL Introduction

Martin Raszyk¹

October 9, 2019

¹using material from Andreas Bärtschi, Petar Ivanov, Chih-Hung Liu, and Daniel Wolleb


Roadmap

- ▶ definition of a graph and its representations
- ▶ declaring and initializing a graph in BGL
- ▶ examples of standard graph algorithms in BGL
- ▶ tutorial problem: from a problem statement to a full solution


Graph: Definition

Definition

A *graph* $G = (V, E)$ consists of a set of n vertices (nodes) V and m edges E .



Undirected graph.



Directed graph.

Directed Graph: Representation

Adjacency matrix:


$$\begin{array}{c} a \\ b \\ c \\ d \end{array} \begin{pmatrix} a & b & c & d \\ 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 \end{pmatrix}$$

Space complexity: $\Theta(n^2)$.

Adjacency list: **USE THIS REPRESENTATION!**

Vertex	List of neighbors
a	[b, c]
b	[c]
c	[d]
d	[a]

Space complexity: $\Theta(n + m)$.




Adjacency List in C++ Standard Library

```
#include <vector>

typedef std::vector<int>      neighbor_list;
typedef std::vector<neighbor_list> cpp_graph;
```

vertex	neighbor_list
0	[1, 2]
1	[2]
2	[3]
3	[0]



Adjacency List in BGL

► C++ Standard Library

```
#include <vector>

typedef std::vector<int>          neighbor_list;
typedef std::vector<neighbor_list> cpp_graph;
```

► BGL


```
#include <boost/graph/adjacency_list.hpp>

typedef boost::adjacency_list<boost::vecS, boost::vecS, boost::directedS> graph;
```

Warm-up: Initialize a Graph...

```
void init_graph()
{
    graph G(4);

    boost::add_edge(0, 1, G);
    boost::add_edge(1, 2, G);
    boost::add_edge(2, 3, G);
    boost::add_edge(3, 0, G);
    boost::add_edge(0, 2, G);
}
```



WARNING!

`boost::add_edge(0, 7, G);` would extend the vertex set of `G` to *eight* vertices!

Warm-up: ...and Iterate over its Edges

► all edges:

```
typedef boost::graph_traits<graph>::edge_iterator edge_it;

edge_it e_beg, e_end;
for (boost::tie(e_beg, e_end) = boost::edges(G); e_beg != e_end; ++e_beg) {
    std::cout << boost::source(*e_beg, G) << " " << boost::target(*e_beg, G) << "\n";
}
```

► neighbors of a vertex:


```
typedef boost::graph_traits<graph>::out_edge_iterator out_edge_it;

out_edge_it oe_beg, oe_end;
for (boost::tie(oe_beg, oe_end) = boost::out_edges(0, G); oe_beg != oe_end; ++oe_beg) {
    assert(boost::source(*oe_beg, G) == 0);
    std::cout << boost::target(*oe_beg, G) << "\n";
}
} /* end of function init_graph */
```



Weighted Graph: Definition

Definition

A *weighted* graph $G = (V, E, w)$ consists of a set of vertices V , edges E , and a *weight function* $w : E \rightarrow \mathbb{Z}$.



Undirected weighted graph.



Directed weighted graph.

Weighted Graph in BGL

Definition

A *weighted* graph $G = (V, E, w)$ consists of a set of vertices V , edges E , and a *weight function* $w : E \rightarrow \mathbb{Z}$.

```
#include <boost/graph/adjacency_list.hpp>

typedef boost::adjacency_list<boost::vecS, boost::vecS, boost::directedS,
    boost::no_property,                    // no vertex property
    boost::property<boost::edge_weight_t, int> // interior edge weight property
> weighted_graph;


typedef boost::property_map<weighted_graph, boost::edge_weight_t>::type weight_map;
```

Initialize a Weighted Graph

```
typedef boost::property_map<weighted_graph, boost::edge_weight_t>::type weight_map;  
typedef boost::graph_traits<weighted_graph>::edge_descriptor edge_desc;
```

```
void init_weighted_graph()  
{  
    weighted_graph G(4);  
    weight_map weights = boost::get(boost::edge_weight, G);  
  
    edge_desc e;  
    e = boost::add_edge(0, 1, G).first; weights[e]=0;  
    e = boost::add_edge(1, 2, G).first; weights[e]=2;  
    e = boost::add_edge(2, 3, G).first; weights[e]=1;  
    e = boost::add_edge(3, 0, G).first; weights[e]=3;  
    e = boost::add_edge(0, 2, G).first; weights[e]=3;  
}
```

- ▶ `edge_weight_t`: a C++ type
- ▶ `edge_weight`: a C++ value




Predefined Vertex and Edge Properties

Some *predefined* vertex and edge properties:


- ▶ `vertex_degree_t`
- ▶ `vertex_distance_t`
- ▶ `edge_weight_t`
- ▶ `edge_capacity_t`
- ▶ `edge_residual_capacity_t`
- ▶ `edge_reverse_t`

All property maps must be initialized and maintained **manually**!

Shortest Path between Two Vertices



Shortest path from 0 to 2.



Longer path from 0 to 2.

Distance between Two Vertices: Dijkstra's Algorithm

```
#include <boost/graph/dijkstra_shortest_paths.hpp>

int dijkstra_dist(const weighted_graph &G, int s, int t) {
    int n = boost::num_vertices(G);
    std::vector<int> dist_map(n);

    boost::dijkstra_shortest_paths(G, s,
        boost::distance_map(boost::make_iterator_property_map(
            dist_map.begin(), boost::get(boost::vertex_index, G))));

    return dist_map[t];
}
```

Time complexity of `boost::dijkstra_shortest_paths` is $O(n \log n + m)$.


Shortest Path between Two Vertices: Dijkstra's Algorithm

```
#include <boost/graph/dijkstra_shortest_paths.hpp>
typedef boost::graph_traits<weighted_graph>::vertex_descriptor vertex_desc;

int dijkstra_path(const weighted_graph &G, int s, int t, std::vector<vertex_desc> &path) {
    int n = boost::num_vertices(G);
    std::vector<int> dist_map(n);
    std::vector<vertex_desc> pred_map(n);

    boost::dijkstra_shortest_paths(G, s,
        boost::distance_map(boost::make_iterator_property_map(
            dist_map.begin(), boost::get(boost::vertex_index, G)))
/* dot! */ .predecessor_map(boost::make_iterator_property_map(
    pred_map.begin(), boost::get(boost::vertex_index, G))));

    int cur = t;
    path.clear(); path.push_back(cur);
    while (s != cur) {
        cur = pred_map[cur];
        path.push_back(cur);
    }
    std::reverse(path.begin(), path.end());
    return dist_map[t];
}
```




Shortest path edges are **red**.
Predecessor edges are **blue**.

Minimum Spanning Tree

Definition

A *minimum spanning tree* of a connected undirected weighted graph $G = (V, E)$ is an acyclic subgraph of G connecting all vertices in V and having the minimum sum of edge weights.



Minimum Spanning Tree: Kruskal's Algorithm

```
#include <boost/graph/kruskal_min_spanning_tree.hpp>

typedef boost::adjacency_list<boost::vecS, boost::vecS, boost::undirectedS,
    boost::no_property, boost::property<boost::edge_weight_t, int> > weighted_graph;
typedef boost::graph_traits<weighted_graph>::edge_descriptor      edge_desc;

void kruskal(const weighted_graph &G) {
    std::vector<edge_desc> mst;    // vector to store MST edges (not a property map!)

    boost::kruskal_minimum_spanning_tree(G, std::back_inserter(mst));

    for (std::vector<edge_desc>::iterator it = mst.begin(); it != mst.end(); ++it) {
        std::cout << boost::source(*it, G) << " " << boost::target(*it, G) << "\n";
    }
}
```


Time complexity of `boost::kruskal_minimum_spanning_tree` is $O(m \log m)$.

Maximum Matching


Definition

A *maximum* matching in an undirected graph $G = (V, E)$ is a subset of its edges with maximum cardinality such that no two edges in the matching share any endpoint.

Remark. A *maximal* matching (i.e., one which cannot be further extended) can be obtained by a simple greedy algorithm. It is *not* necessarily a maximum matching!



Maximal matching.



Maximum matching.



Maximum (perfect) matching.

Maximum Matching: Edmond's Algorithm

```
#include <boost/graph/max_cardinality_matching.hpp>

typedef boost::adjacency_list<boost::vecS, boost::vecS, boost::undirectedS> graph;
typedef boost::graph_traits<graph>::vertex_descriptor vertex_desc;

void maximum_matching(const graph &G) {
    int n = boost::num_vertices(G);
    std::vector<vertex_desc> mate_map(n); // exterior property map
    const vertex_desc NULL_VERTEX = boost::graph_traits<graph>::null_vertex();

    boost::edmonds_maximum_cardinality_matching(G,
        boost::make_iterator_property_map(mate_map.begin(), boost::get(boost::vertex_index, G)));
    int matching_size = boost::matching_size(G,
        boost::make_iterator_property_map(mate_map.begin(), boost::get(boost::vertex_index, G)));

    for (int i = 0; i < n; ++i) {
        // mate_map[i] != NULL_VERTEX: the vertex is matched
        // i < mate_map[i]: visit each edge in the matching only once
        if (mate_map[i] != NULL_VERTEX && i < mate_map[i]) std::cout << i << " " << mate_map[i] << "\n";
    }
}
```

Time complexity of `boost::edmonds_maximum_cardinality_matching` is $O(mn \cdot \alpha(m, n))$.

BGL Outlook – More Algorithms Available!

To be covered in the remainder of this tutorial:

- ▶ breadth-first search (`breadth_first_search`)
- ▶ strongly connected components in directed graphs (`strong_components`)

To be covered in upcoming tutorials:

- ▶ maximum-flow algorithms
- ▶ minimum-cost maximum-flow algorithms


Exercise:

- ▶ connected components in undirected graphs (`connected_components`)
- ▶ check if a graph is bipartite (`is_bipartite`)
- ▶ minimum spanning tree (`kruskal_minimum_spanning_tree`, `prim_minimum_spanning_tree`)


Tutorial Problem: Universal Vertex

Definition


Let $G = (V, E)$ be a directed graph. We call a vertex $u \in V$ *universal* if every other vertex $v \in V$ can be reached from u via a directed path.



Any vertex is universal.



Only 0 is universal.




No vertex is universal.

Tutorial Problem: The Problem Statement

Input A directed unweighted graph $G = (V, E)$.

Output All universal vertices in G .



Tutorial Problem: Checking if a Vertex is Universal

```
#include <boost/graph/breadth_first_search.hpp>
#include <boost/graph/properties.hpp>

typedef boost::adjacency_list<boost::vecS, boost::vecS, boost::directedS> graph;
typedef boost::default_color_type color;
const color black = boost::color_traits<color>::black(); // visited by BFS
const color white = boost::color_traits<color>::white(); // not visited by BFS

bool is_universal(const graph &G, int u) { // Is u universal in G?
    int n = boost::num_vertices(G);
    std::vector<color> vertex_color(n); // exterior property map

    boost::breadth_first_search(G, u,
        boost::color_map(boost::make_iterator_property_map(
            vertex_color.begin(), boost::get(boost::vertex_index, G))));

    // u is universal iff no vertex is white (i.e., all vertices are reachable from u)
    return (std::find(vertex_color.begin(), vertex_color.end(), white) == vertex_color.end());
}
```

Time complexity of `boost::breadth_first_search` is $O(n + m)$.

Tutorial Problem: A Slow Solution


```
void testcase() {  
    int n, m;  
    std::cin >> n >> m;  
    graph G(n);  
  
    for (int i = 0; i < m; ++i) {  
        int u, v;  
        std::cin >> u >> v;  
        boost::add_edge(u, v, G);  
    }  
  
    for (int i = 0; i < n; ++i) {  
        if (is_universal(G, i)) std::cout << i << " ";  
    }  
    std::cout << "\n";  
}
```

Time complexity of testcase is $O(n^2 + nm)$.


Tutorial Problem: Strongly Connected Components

Definition

A *strongly connected component* (SCC) of a directed graph $G = (V, E)$ is any maximal subset of vertices $C \subseteq V$ such that all vertices in C are pairwise reachable (via directed paths).



A directed graph G .



The condensation of G .


Lemma

The condensation of a directed graph is acyclic.


Tutorial Problem: Source SCC

Definition

A strongly connected component is called a *source* if it has no incoming edges from other SCCs.



A directed graph G .



The condensation of G .

Lemma

Any non-empty directed graph has a source SCC.

Lemma

If a directed graph has a single source SCC, then it is a universal vertex in its condensation.

Tutorial Problem: Full Solution (I)

```
#include <boost/graph/adjacency_list.hpp>
#include <boost/graph/strong_components.hpp>

typedef boost::adjacency_list<boost::vecS, boost::vecS, boost::directedS> graph;
typedef boost::graph_traits<graph>::edge_iterator      edge_it;

void testcase() {
    int n, m;
    std::cin >> n >> m;
    graph G(n);

    for (int i = 0; i < m; ++i) {
        int u, v;
        std::cin >> u >> v;
        boost::add_edge(u, v, G);
    }
```

Tutorial Problem: Full Solution (II) — Strongly Connected Components

```
// scc_map[i]: index of SCC containing i-th vertex
std::vector<int> scc_map(n); // exterior property map
// nsc: total number of SCCs
int nsc = boost::strong_components(G,
    boost::make_iterator_property_map(scc_map.begin(), boost::get(boost::vertex_index, G)));
```

Time complexity of `boost::strong_components` is $O(n + m)$.

Tutorial Problem: Full Solution (III) – Source SCCs

```
// is_src[i]: is i-th SCC a source?
std::vector<bool> is_src(nsc, true);
edge_it ebegin, eend;

for (boost::tie(ebegin, eend) = boost::edges(G); ebegin != eend; ++ebegin) {
    int u = boost::source(*ebegin, G), v = boost::target(*ebegin, G);
    // edge (u, v) in G implies that component scc_map[v] is not a source
    if (scc_map[u] != scc_map[v]) is_src[scc_map[v]] = false;
}
```

Tutorial Problem: Full Solution (IV) – Finding All Universal Vertices

```
int src_count = std::count(is_src.begin(), is_src.end(), true);
if (src_count > 1) { // no universal vertex among multiple SCCs
    std::cout << "\n";
    return;
}
assert(src_count == 1); // recall property of the condensation DAG (directed acyclic graph)

// all vertices in the single source SCC are universal
for (int v = 0; v < n; ++v) {
    if (is_src[scc_map[v]]) std::cout << v << " ";
}
std::cout << "\n";
} /* end of function testcase */
```

Time complexity of testcase is $O(n + m)$.

Tutorial Problem: Full Solution (V) – Main Function

```
int main()
{
    int T;
    std::cin >> T;

    while(T-->0) testcase();

    return 0;
}
```